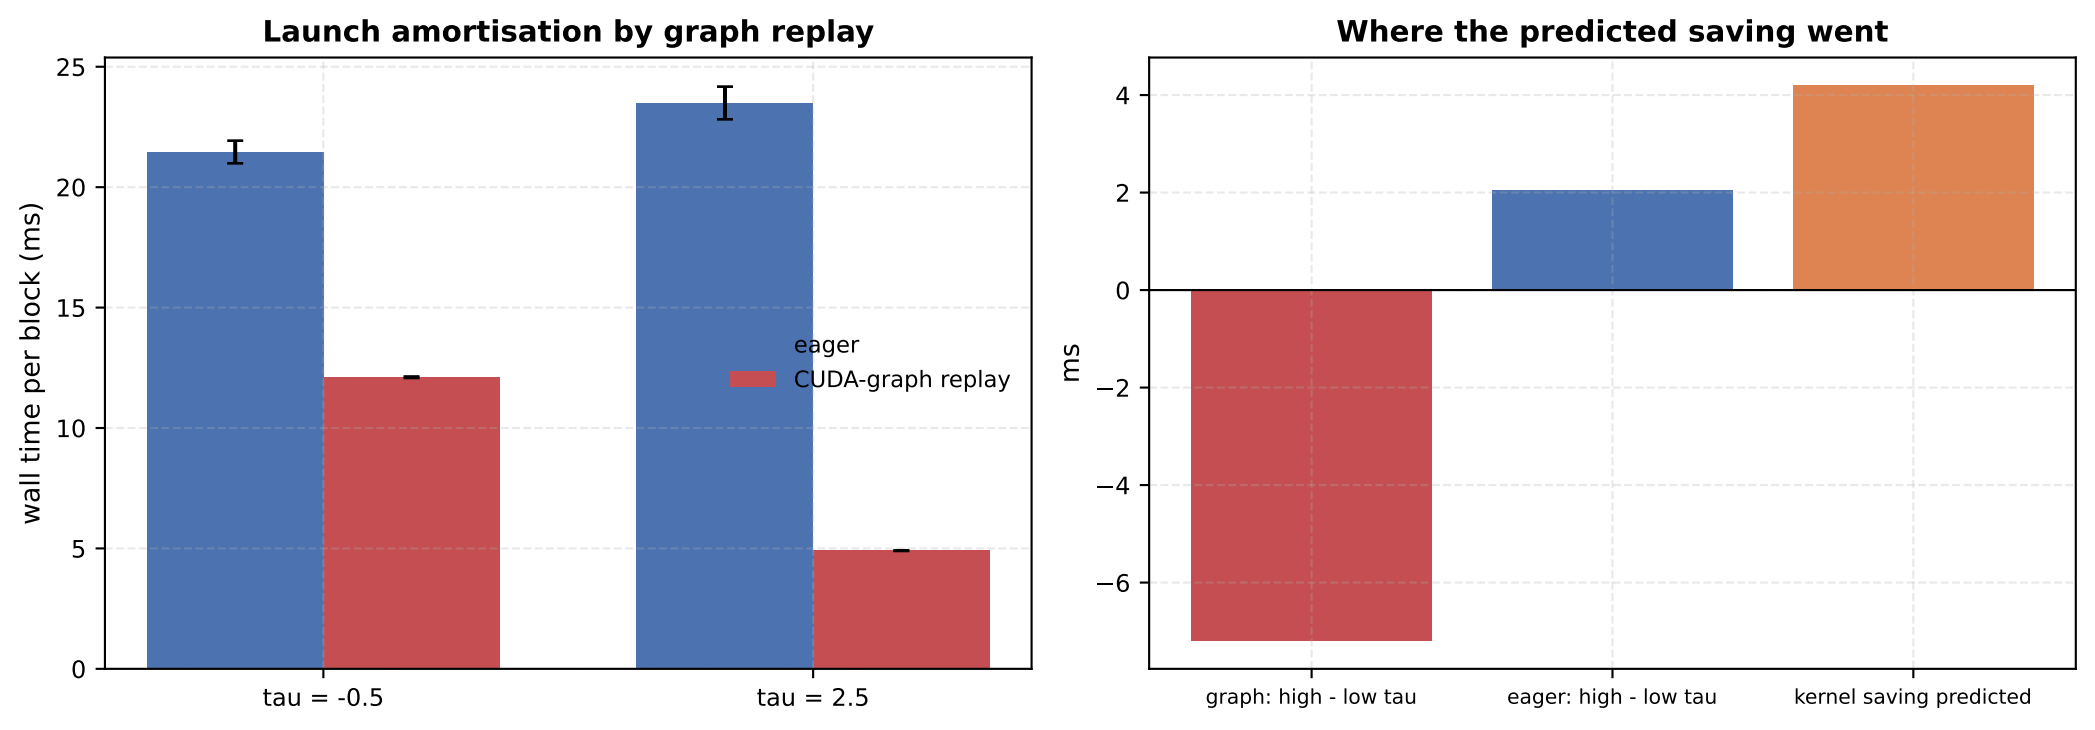
Launch amortisation by graph replay
wall time per block (ms)
25
20
15
10
5
0
eager
CUDA-graph replay
tau = -0.5
tau = 2.5
Where the predicted saving went
ms
4
2
0
−2
−4
−6
graph: high - low tau
eager: high - low tau
kernel saving predicted